

# A fast direct solver for non-oscillatory integral equations

Kenneth L. Ho

Courant Institute, New York University

Argonne Lab, Jan 2012

# Outline

- 1 Integral equations
- 2 Fast iterative solvers
- 3 Fast direct solvers
- 4 Numerical results

# A model problem

Consider Laplace's equation with Dirichlet boundary conditions:

$$\Delta u = 0 \quad \text{in } \Omega \subset \mathbb{R}^3 , \quad u = f \quad \text{on } \partial\Omega.$$

How to find  $u$ ?

# A model problem

Consider Laplace's equation with Dirichlet boundary conditions:

$$\Delta u = 0 \quad \text{in } \Omega \subset \mathbb{R}^3 , \quad u = f \quad \text{on } \partial\Omega.$$

How to find  $u$ ? Use a **double-layer potential** representation

$$u(x) = \int_{\partial\Omega} \frac{\partial G}{\partial \nu(y)}(x, y) \sigma(y) ds(y) \quad \text{in } \Omega,$$

where  $G(x, y) = 1/(4\pi|x - y|)$  is the Laplace Green's function,  $\nu$  is the unit outer surface normal, and  $\sigma$  is an unknown surface density.

# A model problem

Consider Laplace's equation with Dirichlet boundary conditions:

$$\Delta u = 0 \quad \text{in } \Omega \subset \mathbb{R}^3 , \quad u = f \quad \text{on } \partial\Omega.$$

How to find  $u$ ? Use a **double-layer potential** representation

$$u(x) = \int_{\partial\Omega} \frac{\partial G}{\partial \nu(y)}(x, y) \sigma(y) ds(y) \quad \text{in } \Omega,$$

where  $G(x, y) = 1/(4\pi|x - y|)$  is the Laplace Green's function,  $\nu$  is the unit outer surface normal, and  $\sigma$  is an unknown surface density. Then appeal to classical potential theory to derive the **integral equation**

$$-\frac{1}{2}\sigma(x) + \int_{\partial\Omega} \frac{\partial G}{\partial \nu(y)}(x, y) \sigma(y) ds(y) = f(x) \quad \text{on } \partial\Omega.$$

# Why integral equations?

- Well-conditioned (second-kind Fredholm equation)
- Accurate derivatives (differentiate under the integral)
- High-order methods (quadrature)
- Adaptive (often with dimensional reduction)

# Why integral equations?

- Well-conditioned (second-kind Fredholm equation)
- Accurate derivatives (differentiate under the integral)
- High-order methods (quadrature)
- Adaptive (often with dimensional reduction)

Mathematically, the “right” thing to do: try to write down as much of the solution as possible before turning to numerics.

# An exaggerated example

Consider Poisson's equation with free-space boundary conditions:

$$-\Delta u = f \quad \text{in } \Omega \subset \mathbb{R}^3 , \quad u = \mathcal{O}\left(\frac{1}{|x|}\right) \quad \text{as } |x| \rightarrow \infty.$$

# An exaggerated example

Consider Poisson's equation with free-space boundary conditions:

$$-\Delta u = f \quad \text{in } \Omega \subset \mathbb{R}^3 , \quad u = \mathcal{O}\left(\frac{1}{|x|}\right) \quad \text{as } |x| \rightarrow \infty.$$

**Invert** the differential operator:

$$u(x) = [(-\Delta)^{-1} f](x) = \int_{\Omega} G(x, y) f(y) dv(y).$$

In this case, there is no integral equation; the solution is just an integral.

# An exaggerated example

Consider Poisson's equation with free-space boundary conditions:

$$-\Delta u = f \quad \text{in } \Omega \subset \mathbb{R}^3 , \quad u = \mathcal{O}\left(\frac{1}{|x|}\right) \quad \text{as } |x| \rightarrow \infty.$$

**Invert** the differential operator:

$$u(x) = [(-\Delta)^{-1} f](x) = \int_{\Omega} G(x, y) f(y) dv(y).$$

In this case, there is no integral equation; the solution is just an integral.

## Procedure for general boundary conditions

- Integrate volume term
- Solve for boundary correction

# Numerical considerations

Let  $A \in \mathbb{C}^{N \times N}$  be a matrix discretization of some Green's function integral operator. (How to discretize?) Observe that  $A$  is **dense**.

- Cost of applying  $A$ :  $\mathcal{O}(N^2)$
- Cost of inverting  $A$ :  $\mathcal{O}(N^3)$

# Numerical considerations

Let  $A \in \mathbb{C}^{N \times N}$  be a matrix discretization of some Green's function integral operator. (How to discretize?) Observe that  $A$  is **dense**.

- Cost of applying  $A$ :  $\mathcal{O}(N^2)$
- Cost of inverting  $A$ :  $\mathcal{O}(N^3)$

This is contrast to competing methods based on finite differences or finite elements, which yield **sparse** matrices. Historically, this relative expense was a primary cause for the dearth of integral equations in numerical computing (except where there was no choice).

# Fast iterative solvers

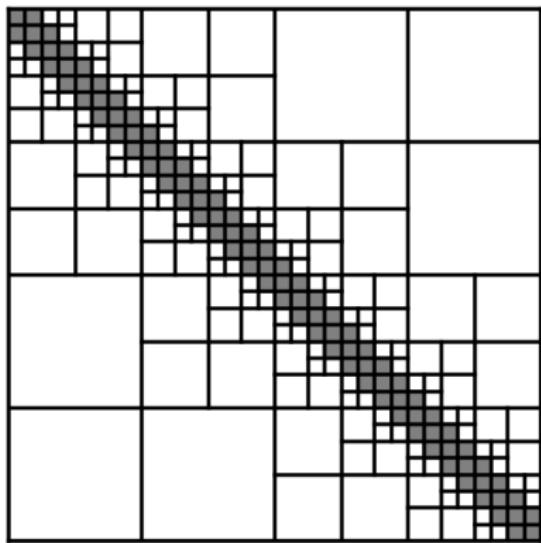
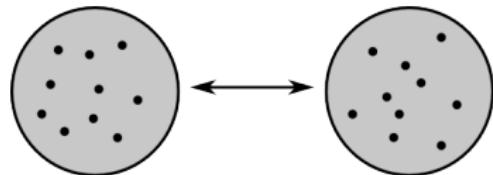
In the 1980s, fast algorithms to apply  $A$  in only  $\mathcal{O}(N)$  or  $\mathcal{O}(N \log N)$  time were developed:

- Treecode (Barnes and Hut, 1986)
- Fast multipole method (Greengard and Rokhlin, 1987)
- Panel clustering (Hackbusch and Nowak, 1989)

Combined with Krylov methods (e.g., GMRES), such techniques enabled fast iterative solution with only  $\mathcal{O}(N \log N)$  work in many situations.

# Fast multipole (and related) methods

- Non-oscillatory Green's functions have **smooth** far fields
- Interactions between well-separated clusters are **low-rank**
- Exploit smoothness with a hierarchical decomposition of space



## Still more work to do

Fast iterative solvers have been very successful, but they remain **inefficient** in certain important regimes:

- When  $A$  is ill-conditioned (e.g., multiphysics, singular geometries)
- When  $Ax = b$  must be solved with many right-hand sides  $b$  or many perturbations of a base matrix  $A$  (e.g., scattering, optimization, design, time marching)

## Still more work to do

Fast iterative solvers have been very successful, but they remain **inefficient** in certain important regimes:

- When  $A$  is ill-conditioned (e.g., multiphysics, singular geometries)
- When  $Ax = b$  must be solved with many right-hand sides  $b$  or many perturbations of a base matrix  $A$  (e.g., scattering, optimization, design, time marching)

One solution: **direct** solvers.

- Robust: insensitive to conditioning, always works
- Fast solves and inverse updates following initial factorization

## Still more work to do

Fast iterative solvers have been very successful, but they remain inefficient in certain important regimes:

- When  $A$  is ill-conditioned (e.g., multiphysics, singular geometries)
- When  $Ax = b$  must be solved with many right-hand sides  $b$  or many perturbations of a base matrix  $A$  (e.g., scattering, optimization, design, time marching)

One solution: direct solvers.

- Robust: insensitive to conditioning, always works
- Fast solves and inverse updates following initial factorization

Can we accelerate direct solvers to the same extent?

# Fast direct solvers

Much active research in recent years:

- $\mathcal{H}$ -matrices (Hackbusch et al., 1999, 2000, 2002)
- Lippmann-Schwinger in 2D (Chen, 2002)
- FMM structure (Pals, 2004)
- BIEs in 2D (Martinsson and Rokhlin, 2005)
- HSS matrices (Chandrasekaran et al., 2006)
- one-level BIE solver in 3D (Greengard et al., 2009)

Current state of the art			
	1D	2D	3D
Precomp	$\mathcal{O}(N)$	$\mathcal{O}(N^{3/2})$	$\mathcal{O}(N^2)$
Solve	$\mathcal{O}(N)$	$\mathcal{O}(N \log N)$	$\mathcal{O}(N^{4/3})$

# Block-separable matrices

## Definition

A block matrix  $A$  is **block-separable** if

$$\underbrace{\begin{bmatrix} \times & \times \\ \times & \times \end{bmatrix}}_{A_{ij}} = \underbrace{\begin{bmatrix} \times \\ \times \end{bmatrix}}_{L_i} \underbrace{\begin{bmatrix} \times \end{bmatrix}}_{S_{ij}} \underbrace{\begin{bmatrix} \times & \times \end{bmatrix}}_{R_j}, \quad i \neq j.$$

# Block-separable matrices

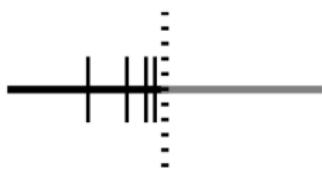
## Definition

A block matrix  $A$  is **block-separable** if

$$\underbrace{\begin{bmatrix} \times & \times \\ \times & \times \end{bmatrix}}_{A_{ij}} = \underbrace{\begin{bmatrix} \times \\ \times \end{bmatrix}}_{L_i} \underbrace{\begin{bmatrix} \times \end{bmatrix}}_{S_{ij}} \underbrace{\begin{bmatrix} \times & \times \end{bmatrix}}_{R_j}, \quad i \neq j.$$

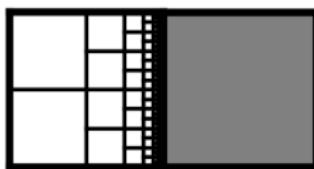
Integral equation matrices are block-separable.

1D



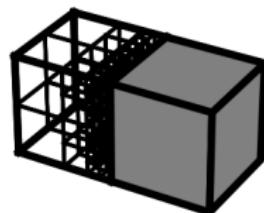
$$\mathcal{O}(\log n)$$

2D



$$\mathcal{O}(n^{1/2})$$

3D



$$\mathcal{O}(n^{2/3})$$

# A fast direct solver for block-separable matrices

If  $A$  is block-separable, then

$$\underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_A = \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_D + \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_L + \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_S + \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_R$$

# A fast direct solver for block-separable matrices

If  $A$  is block-separable, then

$$\underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_A = \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_D + \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_L \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_S \underbrace{\begin{matrix} & & \\ & & \\ & & \end{matrix}}_R$$

and  $A^{-1} = \mathcal{D} + \mathcal{L}\mathcal{S}^{-1}\mathcal{R}$ , where

$$\mathcal{D} = D^{-1} - D^{-1}L\Lambda RD^{-1}, \quad \mathcal{L} = D^{-1}L\Lambda, \quad \mathcal{R} = \Lambda RD^{-1}, \quad \mathcal{S} = \Lambda + S,$$

with  $\Lambda = (RD^{-1}L)^{-1}$ . If  $A$  has  $p \times p$  blocks and each  $S_{ij} \in \mathbb{C}^{k \times k}$ , then  $A^{-1}$  can be computed in  $\mathcal{O}(p(N/p)^3 + (pk)^3)$  operations.

# A sparse matrix perspective

Consider the system

$$Ax = (D + LSR)x = b.$$

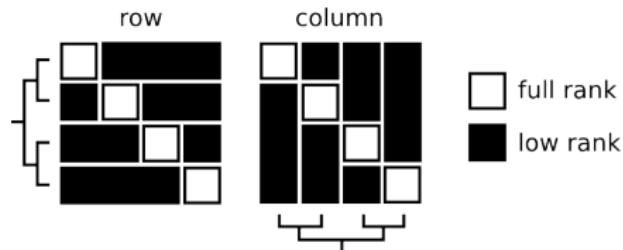
With  $z = Rx$  and  $y = Sz$ , this is equivalent to the **structured sparse** system

$$\begin{bmatrix} D & L \\ R & -I \\ -I & S \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} b \\ 0 \\ 0 \end{bmatrix}.$$

Factor using UMFPACK, SuperLU, MUMPS, Pardiso, etc.

# Hierarchically block-separable matrices

Integral equation matrices are, in fact, **hierarchically block-separable**, i.e., they are block-separable at every level of an octree-type ordering.



In this setting, much more powerful algorithms can be developed.

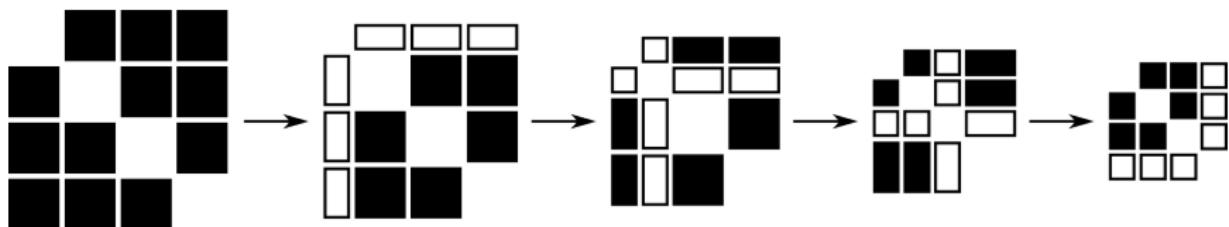
## Definition

An **interpolative decomposition** of a rank- $k$  matrix  $A \in \mathbb{C}^{m \times n}$  is a representation  $A = BP$ , where  $B \in \mathbb{C}^{m \times k}$  is a column-submatrix of  $A$  and  $P \in \mathbb{C}^{k \times n}$  contains the  $k \times k$  identity, such that  $\|P\|$  is small.

- The ID compresses the column space; to compress the row space, apply the ID to  $A^T$ . We call the retained rows and columns **skeletons**.
- Adaptive algorithms exist to compute the ID to any precision  $\epsilon > 0$  in  $\mathcal{O}(kmn)$  time (Cheng et al., 2005; Liberty et al., 2007).

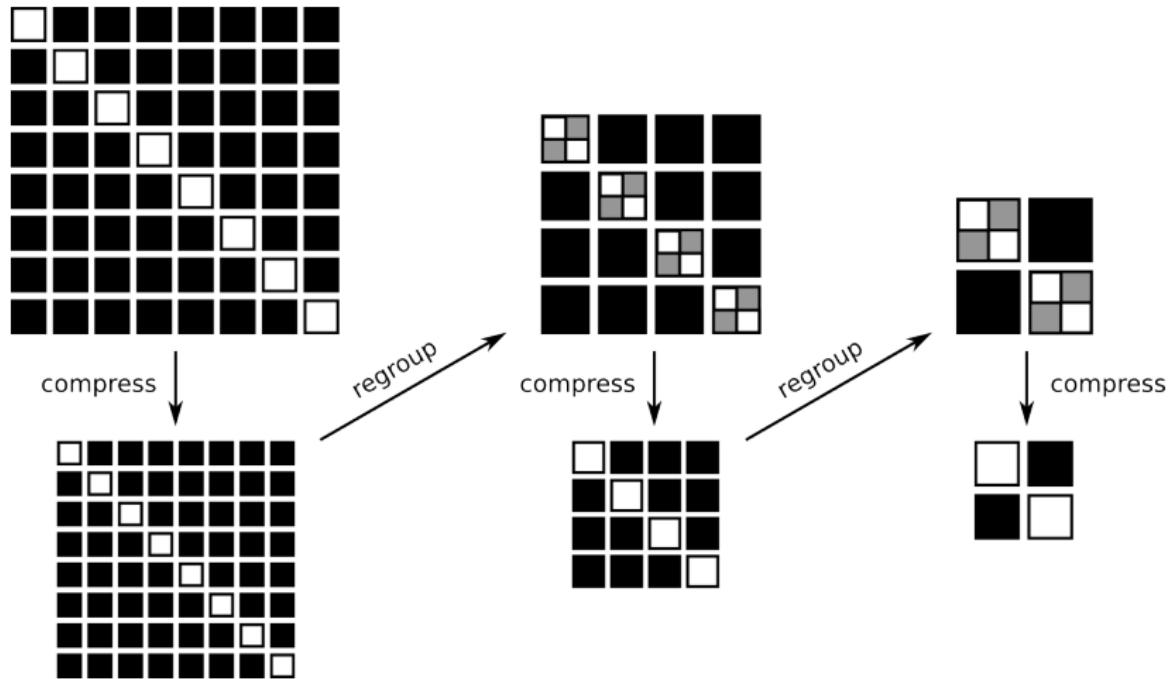
# One-level matrix compression

- ① Compress the row space of each off-diagonal block row with the ID. Let the  $L_i$  be the corresponding row projection matrices.
- ② Compress the column space of each off-diagonal block column with the ID. Let the  $R_j$  be the corresponding column projection matrices.
- ③ Approximate the off-diagonal blocks by  $A_{ij} \approx L_i S_{ij} R_j$  for  $i \neq j$ .



Skeletonization

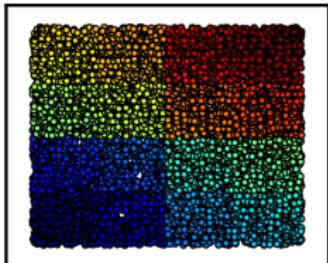
# Multilevel matrix compression



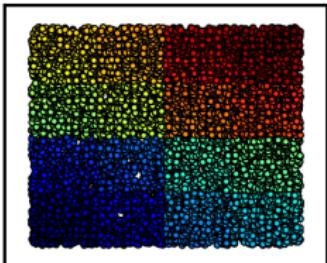
Recursive skeletonization

# Data sparsification

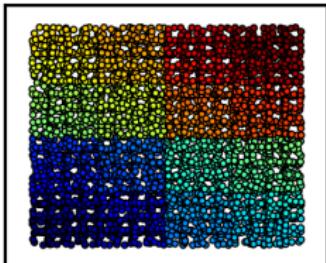
$N_0 = 8192$



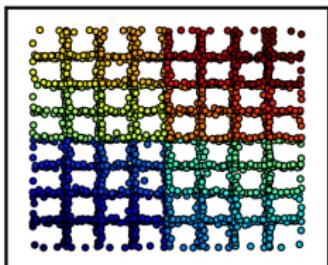
$N_1 = 7134$



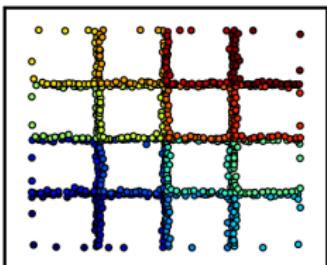
$N_2 = 4138$



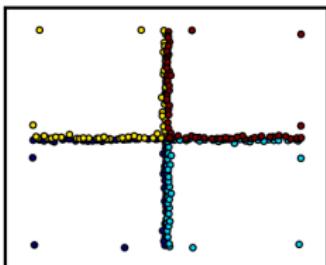
$N_3 = 1849$



$N_4 = 776$



$N_5 = 265$



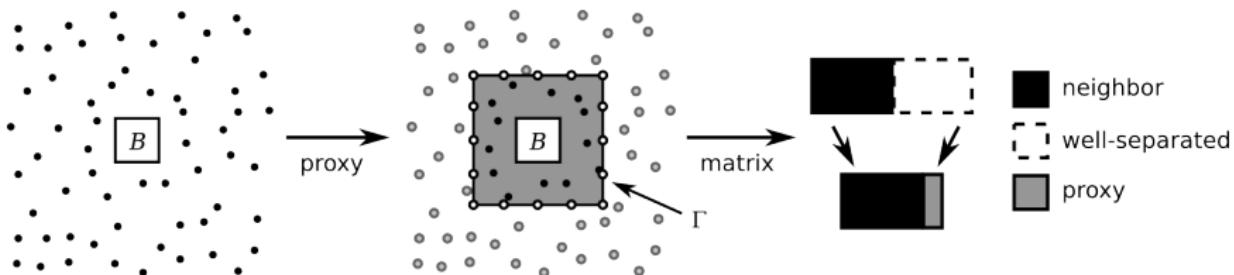
$$G(x, y) = -\frac{1}{2\pi} \log |x - y| , \quad \epsilon = 10^{-3}$$

# Proxy compression

- General compression algorithm is **global** and so at least  $\mathcal{O}(N^2)$
- Use Green's theorem to accelerate:

$$u(x) = \int_{\Gamma} \left[ u(y) \frac{\partial G}{\partial \nu(y)}(x, y) - G(x, y) \frac{\partial u}{\partial \nu(y)}(y) \right] ds(y).$$

- Represent well-separated points with a **local** proxy surface



# Compressed matrix representation

Compressed telescoping matrix representation:

$$A \approx D^{(1)} + L^{(1)} \left[ D^{(2)} + L^{(2)} \left( \dots D^{(\lambda)} + L^{(\lambda)} S R^{(\lambda)} \dots \right) R^{(2)} \right] R^{(1)}$$

# Compressed matrix representation

Compressed telescoping matrix representation:

$$A \approx D^{(1)} + L^{(1)} \left[ D^{(2)} + L^{(2)} \left( \dots D^{(\lambda)} + L^{(\lambda)} S R^{(\lambda)} \dots \right) R^{(2)} \right] R^{(1)}$$

- Efficient storage (**data-sparse**)

Storage (2D volume,  $\epsilon = 10^{-3}$ )

$N$	Uncompressed	Compressed
8192	537 MB	9.65 MB
131072	137 GB	184 MB

# Compressed matrix representation

Compressed telescoping matrix representation:

$$A \approx D^{(1)} + L^{(1)} \left[ D^{(2)} + L^{(2)} \left( \dots D^{(\lambda)} + L^{(\lambda)} S R^{(\lambda)} \dots \right) R^{(2)} \right] R^{(1)}$$

- Efficient storage (data-sparse)
- Fast matrix-vector multiplication (generalized FMM)
- **Fast matrix factorization and inverse application**

Storage (2D volume,  $\epsilon = 10^{-3}$ )

$N$	Uncompressed	Compressed
8192	537 MB	9.65 MB
131072	137 GB	184 MB

# Sparse inverse embedding

Recursively expand in sparse form:

$$\begin{bmatrix} D^{(1)} & L^{(1)} \\ R^{(1)} & -I \\ -I & D^{(2)} & L^{(2)} \\ R^{(2)} & \ddots & \ddots \\ \ddots & D^{(\lambda)} & L^{(\lambda)} \\ R^{(\lambda)} & -I & -I \\ -I & S \end{bmatrix} \begin{bmatrix} x \\ y^{(1)} \\ z^{(1)} \\ \vdots \\ \vdots \\ y^{(\lambda)} \\ z^{(\lambda)} \end{bmatrix} = \begin{bmatrix} b \\ 0 \\ 0 \\ \vdots \\ \vdots \\ 0 \\ 0 \end{bmatrix}$$

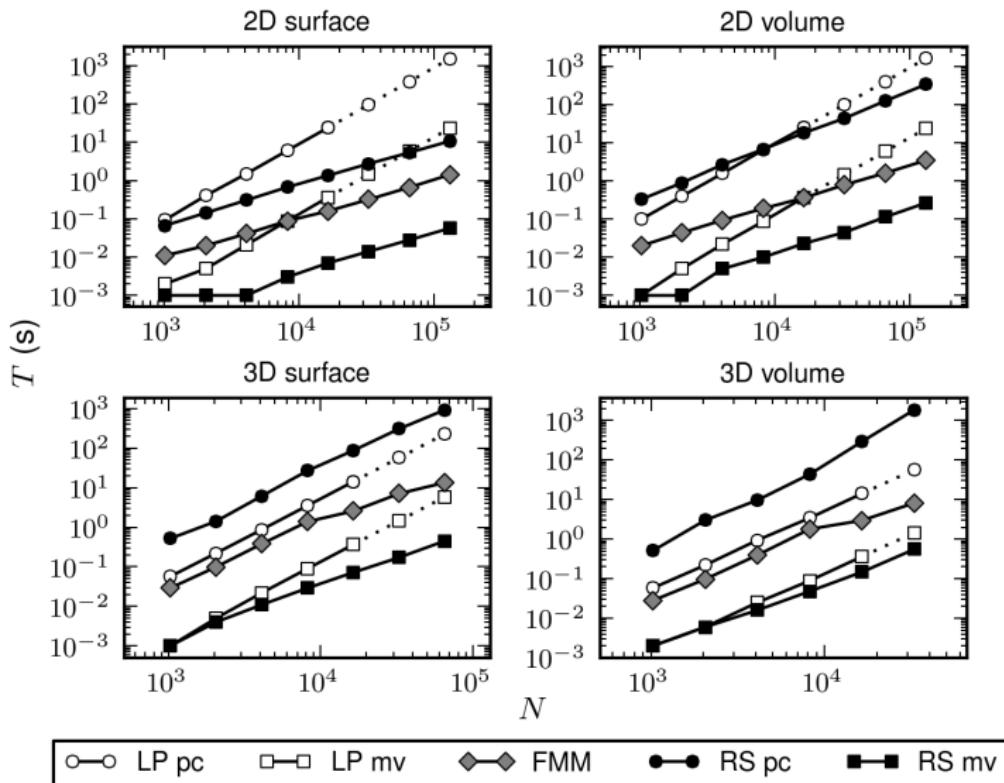
Multilevel **inversion** formula:

$$A^{-1} \approx \mathcal{D}^{(1)} + \mathcal{L}^{(1)} \left[ \mathcal{D}^{(2)} + \mathcal{L}^{(2)} \left( \dots \mathcal{D}^{(\lambda)} + \mathcal{L}^{(\lambda)} \mathcal{S}^{-1} \mathcal{R}^{(\lambda)} \dots \right) \mathcal{R}^{(2)} \right] \mathcal{R}^{(1)}.$$

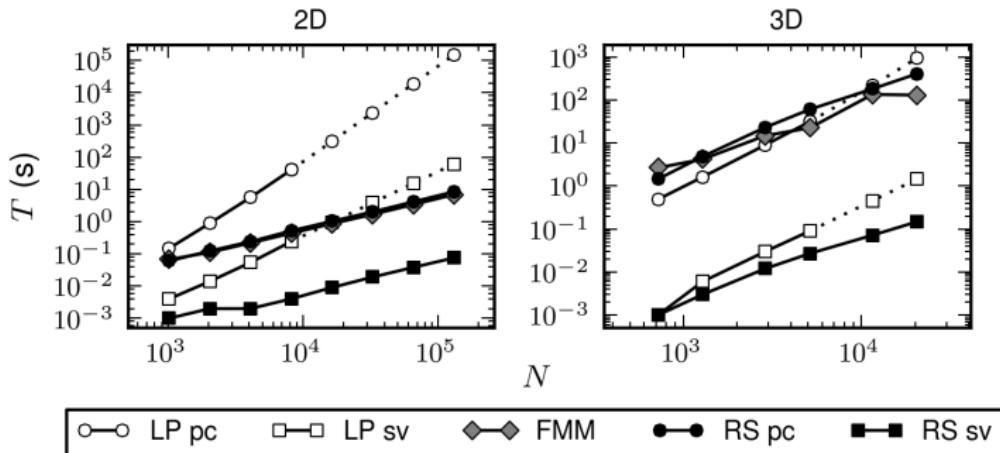
## Some comments

- Mild assumptions: low-rank off-diagonal blocks, Green's theorem
- Based on numerical linear algebra rather than analytic expansions
- **Kernel-independent** (non-oscillatory, elliptic): Laplace, Stokes, Yukawa (screened Poisson), low-frequency Helmholtz, etc.
- Compressed ranks are **optimal** for the problem at hand
- Trade accuracy for speed: user-specified precision
- Naturally parallelizable via block-sweep structure
- Current limitation: optimal complexity only in **1D** (e.g., BIEs in 2D, axisymmetric BIEs in 3D)
- Same ideas can produce an  $\mathcal{O}(N)$  compression-based FMM
- Can also work for PDE formulations (Xia et al., 2009)

# Laplace FMM



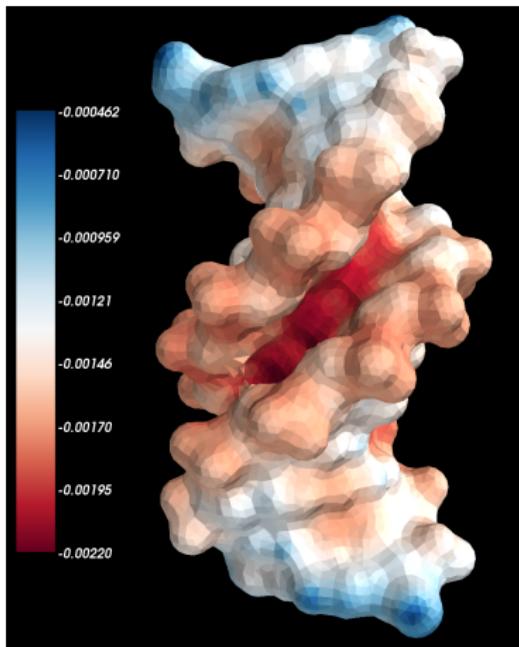
# Laplace BIE solver



- Less memory-efficient than FMM/GMRES
- Solve times are **extremely** fast (in elements/sec)

$\epsilon$	$10^{-3}$	$10^{-6}$	$10^{-9}$
2D	$3.3 \times 10^6$	$2.0 \times 10^6$	$1.7 \times 10^6$
3D	$6.0 \times 10^5$	$1.4 \times 10^5$	$6.2 \times 10^4$

# Molecular electrostatics

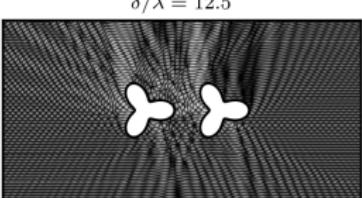
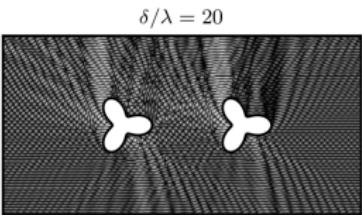
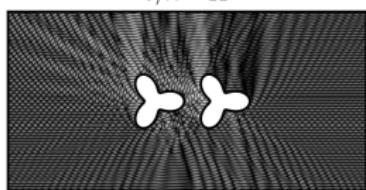
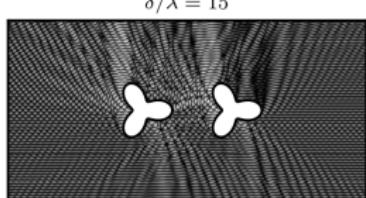
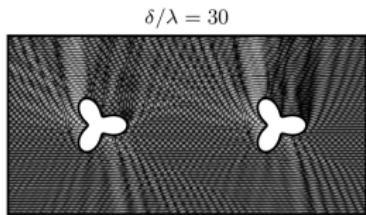


- Piecewise constant dielectric Poisson with interior sources
- $N = 19752, \epsilon = 10^{-3}$
- FMM/GMRES: 27 s
- RS precomp: 578 s
- RS solve: 0.08 s
- Break-even: 25 solves

# Helmholtz problems

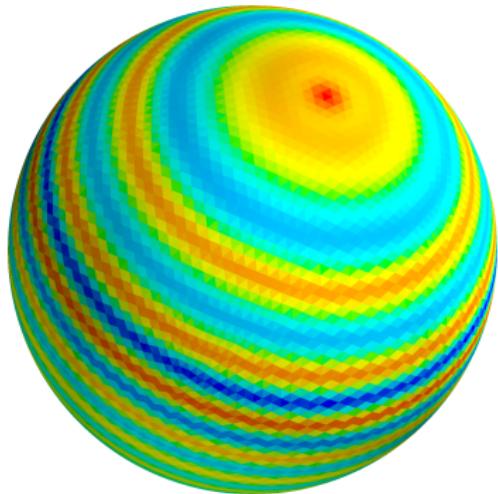
- Interactions are **full-rank** at high frequency
- Asymptotically no acceleration (becomes  $\mathcal{O}(N^3)$  scheme)
- However, remains surprisingly viable:  $200\lambda$  in 2D,  $10\lambda$  in 3D
- At low to moderate frequency, sometimes **superior** to FMM/GMRES due to high iteration counts
- Memory requirements can be a concern
- Effective **preconditioners** for iterative solvers

# Multiple scattering



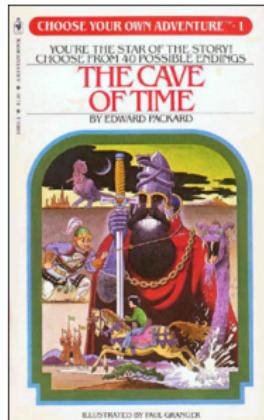
- Each object:  $10\lambda$
- Block inverse preconditioner
- Unpreconditioned: 700 iterations
- Preconditioned: 10 iterations
- $50\times$  speedup

# Approximate inverse preconditioning



- $N = 20480, 10\lambda$
- Precondition with **low-precision** inverse ( $\epsilon = 10^{-3}$ )
- Iterate for full accuracy ( $\epsilon = 10^{-12}$ )
- Unpreconditioned: 190 iterations
- Preconditioned: **6** iterations
- $10 \times$  speedup

# Choose your own adventure

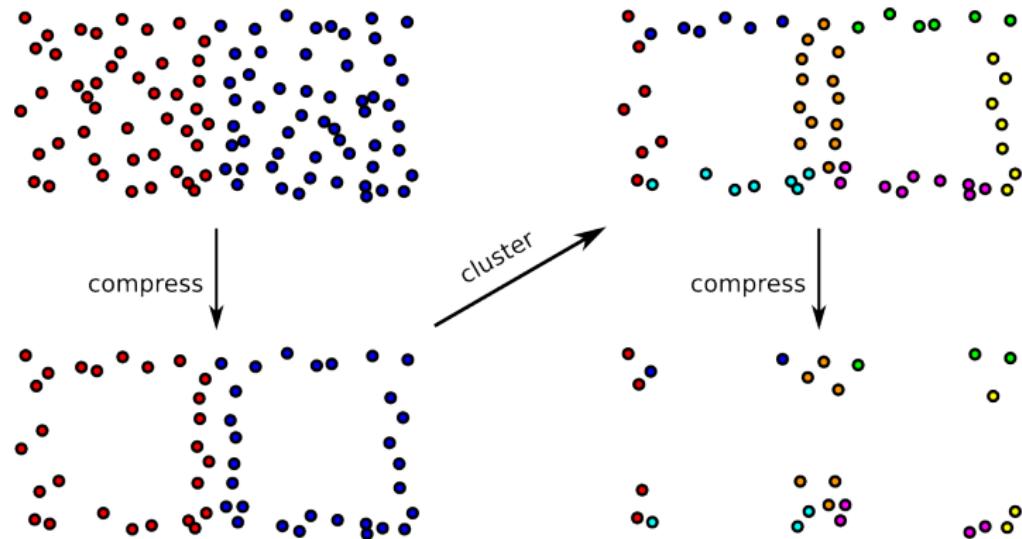


- ① Towards a near-optimal direct solver [▶ Go](#)
- ② Towards a direct solver for oscillatory kernels [▶ Go](#)
- ③ Concluding remarks [▶ Go](#)

# Towards a near-optimal direct solver

- In 2D and above, the number of skeletons grows with the box size
- Need more compression: **skeletonize** the skeletons
- Recall that skeletons line up along box boundaries
- Simple idea: **cluster** according to boundaries and re-compress
- Multiple rounds of compression at each level
- Recurse down on dimensionality: reduce to 1D case
- Expect  $\mathcal{O}(N)$  or  $\mathcal{O}(N \log N)$  complexity

# Recompression schematic



◀ Return

# Towards a direct solver for oscillatory kernels

Previous work:

- BIEs on elongated objects in 2D (Michielssen et al., 1996; Martinsson and Rokhlin, 2007)
- Lippmann-Schwinger in 2D (Chen, 2002)
- BIEs on quasi-planar objects in 3D (Winebrand and Boag, 2009)

No results yet for **general** BIEs!

# Towards a direct solver for oscillatory kernels

Previous work:

- BIEs on elongated objects in 2D (Michielssen et al., 1996; Martinsson and Rokhlin, 2007)
- Lippmann-Schwinger in 2D (Chen, 2002)
- BIEs on quasi-planar objects in 3D (Winebrand and Boag, 2009)

No results yet for **general** BIEs!

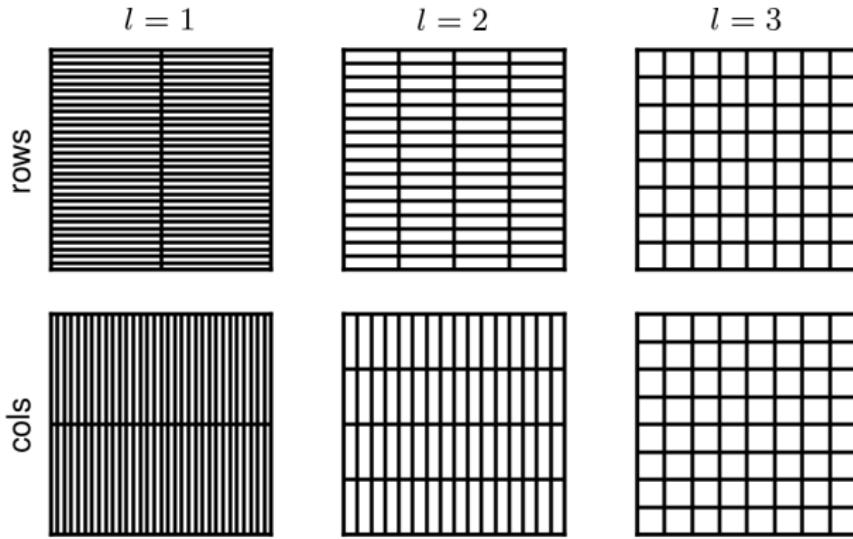
A key observation: at high frequency, the interaction rank between:

- two large boxes is large
- one large box and one small box is **small**

Exploit with the **butterfly algorithm** (Michielssen and Boag, 1996).

# Butterfly algorithm

- Blocks of constant “area” have constant rank
- Opposite tree traversal
- Assume no singularities: e.g., consider Fourier integral operators  
(Candès et al., 2009; Ying, 2009; O’Neil et al., 2010)



# Inverting the butterfly

At each level:

$$A = \begin{bmatrix} L_1^{(1)} S_{11} R_1^{(1)} & L_1^{(2)} S_{12} R_2^{(1)} \\ L_2^{(1)} S_{21} R_1^{(2)} & L_2^{(2)} S_{22} R_2^{(2)} \end{bmatrix}$$

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_1^{(1)} & & L_1^{(2)} & \\ & L_2^{(1)} & & L_2^{(2)} \\ & & L_2^{(1)} & \\ & & & L_2^{(2)} \end{bmatrix} \begin{bmatrix} S_{11} & & & \\ & S_{12} & & \\ & & S_{21} & \\ & & & S_{22} \end{bmatrix} \begin{bmatrix} R_1^{(1)} & & \\ & R_2^{(1)} & \\ R_1^{(2)} & & \\ & R_2^{(2)} & \end{bmatrix}$$

$A$        $L$        $S$        $R$

**Inverse:**  $A^{-1} = R^{-1}S^{-1}L^{-1}$ , or use sparse embedding without  $D$  matrix

# Butterfly comments

- Multilevel and multidimensional extensions are straightforward
- Preliminary experiments suggest fast inversion
- Expect  $\mathcal{O}(N \log N)$  for Fourier integral operators (proxy also available)
- As yet unclear how to treat Helmholtz case (diagonal extraction?)
- Fame, fortune, and glory: revolutionize electromagnetics and imaging

◀ Return

# Concluding remarks

- Fast direct solver for non-oscillatory integral equations (Poisson, Stokes, low-frequency Helmholtz, etc.)
- Based on multilevel matrix compression: pure numerical linear algebra
- Following precomputation, solve times are **very** fast
- Applications: optimization, design, evolution of time-dependent processes in fixed geometries, preconditioning
- Further work: more efficient algorithms, direct solvers for oscillatory kernels, other matrix factorizations